

Design Document: MIT HackTrack

George Ezenna, Stancellous Matoreva, Kairat Ashim, Eric Manzi, Favyen Bastani

Overview

MIT HackTrack will be a hacker community where members of the MIT community share their latest projects, get feedback on their ideas, and find interesting projects to contribute to.

Currently, students don't have an effective medium for discussing projects. Although people frequently communicate project ideas over e-mail, Zephyr, or even social networks like Facebook, these platforms aren't tailored for the task. Long-lived project discussions over e-mail and Zephyr are difficult due to the lack of support for different discussion topics; it's also impossible for new users to join in the middle of a discussion since the conversation history won't be visible. Facebook presents a more viable platform: a student can create a Facebook group for their projects and post development updates, demo links, and screenshots in the group. However, Facebook groups still make it hard for other students to find interesting projects and filter projects by their interests; it lacks a browsable project directory. At the same time, Facebook's Orwellian privacy policy drives away many MIT students.

At the same time, students frequently start working on a large project, but shortly find that it's difficult to make progress by themselves in the small amount of free time that MIT students have. HackTrack enables students to share their idea and quickly see if developers in the MIT hacker community are interested in joining in.

Similarly, getting honest feedback on a project is not straightforward. By posting ideas to hackers that the author doesn't necessarily know already, this is mitigated. Simultaneously, opening up feedback to the community allows students to get suggestions from a more diverse group.

Lastly, it's difficult to keep up-to-date with projects: a student might be interested in several projects that are being developed concurrently, and may not be able to keep track of all of the activity.

Concepts

Asset: an image, video, URL, or other media item relevant to the project. Assets allow users to easily understand the main ideas behind a project when they visit the project's dashboard page. A project owner can upload an image/video file or enter a URL to add an asset. These will then be displayed near the top of the project dashboard for other users to see.

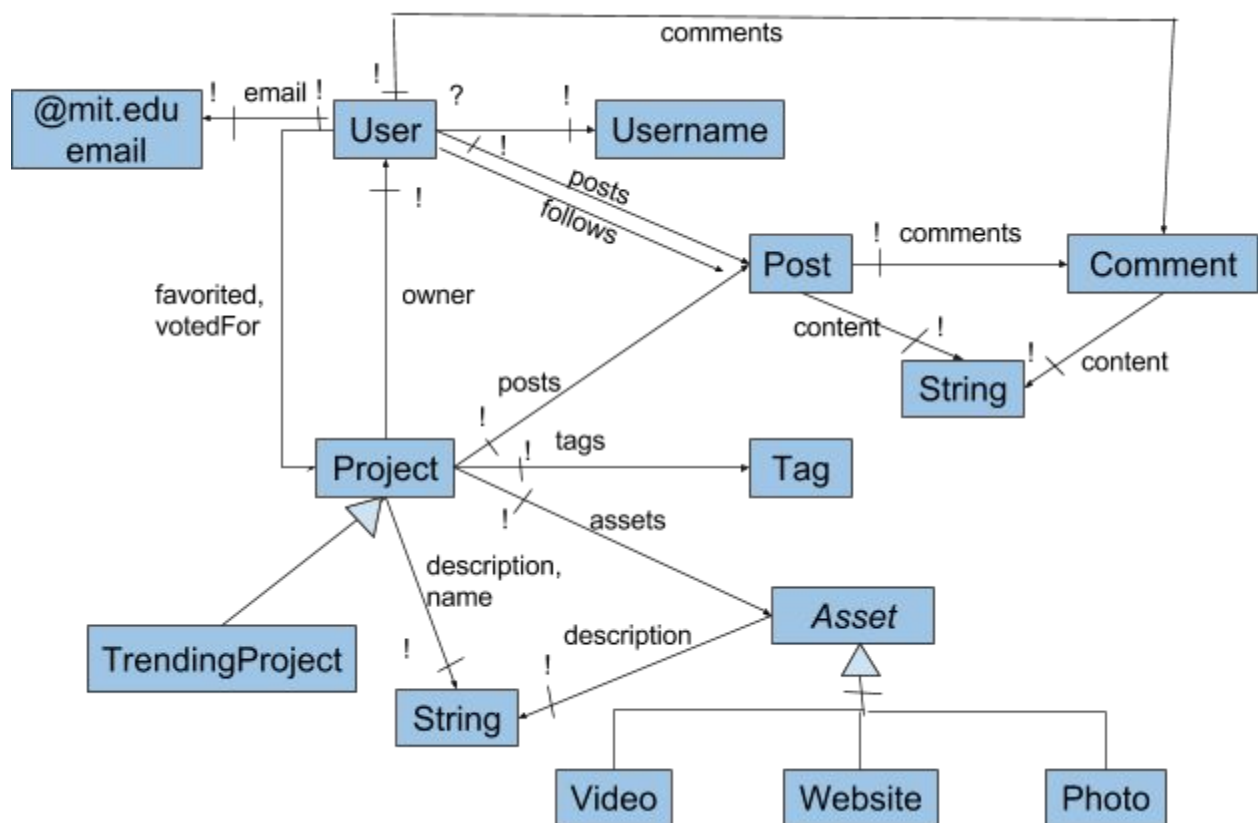
Tag: a key category used to describe the type of the project. Tags make it easy to share and find projects that match a user's interests. Suppose a user has a project consisting of a website where MIT students can share their project ideas. The user may associate the "community" tag with the project to label the project as one that aims to support a particular community. Then, another student can search for the "community" tag and find the project.

Post: a topic of discussion on a project. Posts enable users to provide feedback for a project, and to discuss other project issues. When a user finds a project and has some feedback, the user may make a post with the feedback. Other users can then visit the project and see the post.

Favorites: a set of projects that a user is interested in. Favorites allow users to easily keep up-to-date with projects when there is new activity in the project. A user selects Favorite on the project dashboard to add the project to the user's set of favorites. Then, new project activity will be visible on the user's Activity Feed.

Upvote: a count of how many users have upvoted a project. Upvotes aimed at increasing visibility or providing a rating system for projects that people find interesting. When a user reads about a project and finds it interesting, they upvote it. Projects with the highest votes appear in "Trending projects"

Data Model



Textual constraints:

- Users cannot upvote their own projects
- If a user makes a post, they automatically follow it

Insights and explanations:

- Users can have multiple posts within the same project
- We need to restrict email accounts to single users. Multiple users should not share an email. Since one can create multiple mailing lists, however, we can't ensure that users do not have multiple accounts.
- Posts and comments may be empty.
- Users can choose whether or not to favorite their own projects.

Security Concerns

Security concerns for HackTrack are largely the same as those for standard web applications.

HackTrack will be deployed on a platform-as-a-service (PaaS) provider like Heroku. We assume the provider is trusted, and we additionally consider any platform security issues (e.g. network-level denial of service attacks, physical datacenter security, remote access security) to be the responsibility of the provider (and assume the provider will fulfill that responsibility).

Furthermore, we use HTTPS, and assume that the PaaS provider will ensure that encryption software is up-to-date, so that attackers cannot spoof our service or intercept communications.

We also ignore any physical security on the user-side, where the attacker spies on the user directly.

Thus, in our threat model, the attacker only has the ability to make arbitrary HTTP requests to our application. We then address these web application security concerns:

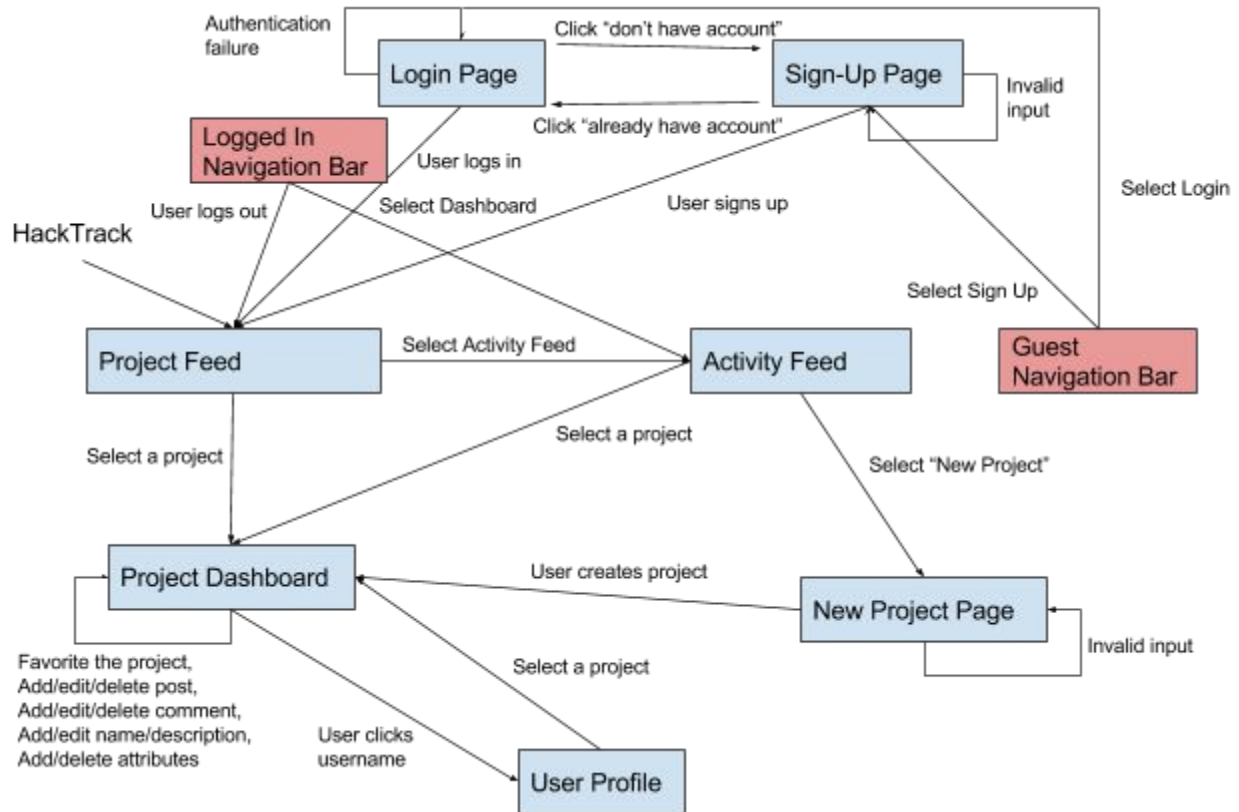
- Cross-site-scripting (XSS) attacks: we use a templating engine that escapes user content when the HTML page is generated. We respond with JSON to asynchronous requests, and ensure that any user content is sanitized in the JSON content on the server.
- Cross-site request forgery (CSRF) attacks: we include a token, stored in the session, to any POST, PUT, and DELETE requests made by the client, and validate the token in the application. We further ensure that GET requests do not affect the server state.
- Database injection attacks: we interact with MongoDB via the Mongoose library; user content is represented as variables in Mongoose model instances, and Mongoose handles data escaping when communicating with the database.
- Access control issues: to avoid accidental cases where access control is missing from some actions, we use routers so that any action that requires user authentication will be placed under a router that includes an authentication middleware.
- Spamming: an attacker may make a large number of sign-up requests with various legitimate e-mail address to cause spam complaints; or, users may simply sign-up with someone else's e-mail address, causing that person to receive

notifications from our system. We rate-limit sign-ups per IP address, and require e-mail address verification before the activating accounts.

- Username/password guessing: we rate limit login attempts based on IP address to mitigate username enumeration or password guessing.
- Database security: our database may become compromised if our PaaS provider is attacked, or we have a copy of the database in a backup system or on a development machine. We want to mitigate the fallout from an event where an attacker obtains a copy of the database. To do so, we hash passwords using a password hashing algorithm like PBKDF2 before storing the hash and salt in the database.
- Password reset form vulnerabilities: a password reset form is needed in case users forget passwords. To prevent attackers from spamming users by repeatedly issuing password reset requests, we rate limit password reset requests to one per day per username, and also require both username and e-mail address. To prevent attackers from resetting the password, we require e-mail address verification and rely on the security of SMTP infrastructure.
- Session management: we use the express-session library to handle sessions and trust the library.
- Outdated components: we subscribe to security updates for node.js modules used in our application, and update the modules routinely.

User Interface


Transition Diagram



Guest Navigation Bar



Logged In Navigation Bar



Login Page

Log In

Don't have an account yet? [Sign up.](#)

Sign Up Page

Sign Up

By signing up, you agree to the [Terms of Service](#).

Already have an account? [Login](#).

Project Feed

- All
- Tag1
- Tag2
- Tag3
- Tag4
- Tag5



Project 1

A brief description of project 1. It is a very interesting project.



9001

Project 2

A brief description of project 2. It is an interesting project.



999

Rockets

Rockets are cool, let's build some!



777

HackTrack

MIT project community. It's for 6.170.



555

Activity Feed

Latest Activity



Ryan Hoover posted [Trumpiñata](#)

3 hours ago



Ryan Hoover posted [59 Illustrated National Parks](#)

21 hours ago



Ben Tossell posted [Tuff: A Novel](#)

a day ago



Ben Tossell posted [The Black Jacobins](#)

a day ago



Ryan Hoover posted [Product Hunt 2.0](#)

a day ago



Ben Tossell posted [Social Media Calendar by Buffer](#)

a day ago



Ben Tossell posted [Launching a Startup in the Digital Age](#)

2 days ago



Ben Tossell posted [99% Invisible - 186: War and Pizza](#)

2 days ago

Project Dashboard



Attributes

Code: <https://github.com/rocketsarecool/controller>

Mailing list: rocketsarecool@mit.edu



Posts

user567

Hi! When are we launching the rockets?

rocketexpert9 (owner)

Tomorrow at 9pm!

user965

That is a good question.

↓ 3 more comments

↓ 23 more posts

New Project Page

Create a New Project

Name

Description


tag1 X

tag2 X

Add a tag...

Create Project

User Profile



alyssaph

Joined on 1 January 2016

Name: Alyssa P. Hacker

Interests: web security,
cryptography

Active Projects

Hacking the NSA

Looking for vulnerabilities in nsa.gov (and helping them fix the issues)!

↑

23

TorDeploy

Easily deploy web applications on Tor.

↑

17

More Projects

AutoAudit

Automatically find XSS and CSRF vulnerabilities in node.js apps.

↑

764

Design Challenges

1. How to authenticate users as members of the MIT community?

Since we are limiting the use of this app to members of the MIT community, we think it's necessary to authenticate users as MIT members when they sign up. We considered the following options:

A. MIT certificates

- **Pro:** Authentication by MIT certificates would be a robust authentication method which is guaranteed to work so long as the user has a valid MIT certificate.
- **Con:** However, to implement MIT's certificate authentication, one would need an SSL packet and approval from IS&T for a signed server certificate.

B. Nodemailer

- **Pro:** Relatively easy to use and does not modify the User model. This method of authentication asks the user for a valid MIT email and sends an email with a verification link to the user's email. Clicking this link would then authenticate the user as an MIT student
- **Con:** Need to store users who have not yet verified their accounts and delete them after a certain deadline if they are not verified and otherwise move them to the user collection when their accounts are verified.

Team verdict: In the interest of saving time, we decided to go with the latter option, Nodemailer. We decided that the time it would take to get approval from IS&T could be unboundedly long. Nodemailer avoids dependence on external factors. To sign up, a user must enter their MIT email, a username and a password. The email they entered is then parsed to verify that it is an MIT email (alum email also allowed.) If it passes, an activation link is then sent to that email and the user must click on that link within 24 hours to be registered to our app. This ensures that the MIT email they entered is in fact their own.

2. Best way to present updates on projects users are interested in

Our options:

- A. If a user likes a project, they favorite it, which adds it to their list of favorite projects. New activity on that project will be shown in that user's activity feed. Additionally, users follow other users, which adds all of the followee's projects to that user's favorites.

- **Pros:** If user A likes user B's projects, they don't have to individually favorite all of that user B's projects and if user B creates a new project that user A misses in projects feed, then they still get notified of that new project.
 - **Con:** Following users might form a conflicting mental model with favoriting projects and users might find it confusing. For instance, if user A follows user B, that implies that user A favorited user B's projects but if user A unfavorites one of user B's projects, it's not clear whether user A no longer follows user B.
- B. If a user likes a project, they favorite it, which adds it to their list of favorite projects. Recent activity on that project appears in that user's Activity feed. Additionally, users follow posts, which notifies the user when a new comment is added to that post.
- **Pros:** Clear separation between following posts and favoriting projects. Following and unfollowing posts has no effect or relation with favoriting projects. Following posts allows users to be updated about conversations that they are interested in, while favoriting projects adds the project's activity to the user's activity feed.
 - **Con:** Users not able to follow other users

We decided on option B because it greatly improves usability in terms of creating a clear mental model for the user. The inability to follow other users is mitigated by implementing profile pages for users which shows all of their projects.

3. Where to store votes in model

Users vote on projects that they like. In designing the model, we considered the following options:

- A. Have the User model store a list of all the posts the user has upvoted.
- B. Have the Project model store all the users that upvoted it
- C. Have the User store a list of their upvoted projects and have the Project only store its vote count

After some deliberation, we decided to go with the last option, C: The advantage this has over option A is that to get the vote count of a Project, we don't have to query through all the users. It also outdoes option B in having a lower number of queries made whenever a user is upvoting a project: With B, when a user upvotes a project, we would need to search through all the users that have upvoted it and check that the user trying to upvote

the post hasn't upvoted it before. With C however, we only need to check that this post is among the posts the user has upvoted, which is a faster query.

4. Order in which to display projects

One main design challenge of this app is how to display projects to the users on the project feed. If projects are displayed in a sorted manner, i.e. by date and/or votes, some projects will never get their time in the sun. For instance, old projects with low votes will get limited visibility. We considered two possible solutions:

- A. Sort all the projects by date (most recent first)
 - **Pros:** New projects are most visible. Projects with low votes still get shown at the top of the feed
 - **Cons:** Old projects get limited visibility
- B. Sort all the projects by votes.
 - **Pros:** Most voted on project are the most visible. Old projects still get shown at the top of the feed
 - **Cons:** Low vote projects get limited visibility. This is slightly better than A
- C. Projects are sorted by date (most recent first.) Additionally have a "Trending projects" page which shows x posts with the highest vote count, sorted by votes. Also have a "Featured projects" page which shows x posts with the lowest vote count, reverse sorted by date so oldest posts are shown first.
 - **Pros:** This method ensures the most fair visibility to all projects.
 - **Cons:** No incentive to visit "Featured projects" page. Implementation is more time consuming
- D. Compound sort projects by date then by vote count such that all posts from one day are shown before posts from a previous day and within each day, posts are sorted by votes.
 - **Pros:** New projects are most visible. Projects with low votes still get shown at the top of the feed
 - **Cons:** Old projects get limited visibility

We went with option C because it performs the best in distributing visibility among the projects despite the fact that it will take the longest time to implement among the other options. For the MVP, we will not be implementing the "Featured projects" since the app's core functionality does not require it.

5. Update notification: email or in-app

If a user follows a post, they will receive updates when someone comments on or upvotes that post. Potential solutions include:

- A. Users receive email whenever a post they follow is updated
 - **Pro:** simple to implement
 - **Con:** emails get annoying if a post you followed gets lots of publicity
- B. Notification bar in User's dashboard
 - **Pro:** Less annoying than email
 - **Con:** Implementing a UI-friendly notification bar is more time consuming than it is an improvement over email notification

For the purpose of implementing a useful minimum viable product in time, we chose to use email as our form of notification since implementing a notification bar might turn out to be a huge time drain.

6. Project owners vs Collaborators

Projects typically have more than one person involved. If someone is working on a project with other people and then posts the project to our app, should the collaborators on that project have edit permissions on that project's page? How should the poster attribute credit to his teammates? We considered the following options:

- A. Users add collaborators to a project which gives them the same access rights to a project's description, features... as the original poster.
 - **Pros:** All collaborators receive credit. All are notified about comments
 - **Cons:** Potential asynchronous errors. If two collaborators are editing the same project at the same time, it is not defined whose edits will take effect.
- B. Users optionally write down names of collaborators. The app doesn't check if that user exists in the system.
 - **Pros:** All collaborators still receive credit.
 - **Cons:** Collaborators not notified about feedback on the project.

After some deliberation, we decided to go with option B. There's only one project owner. This greatly simplifies implementation while still attributing credit to a project's collaborators. We imagine that only in rare cases will all the collaborators on a project be signed up for the app which means implementing 'collaborators' would have taken a considerable amount of time while it would only be useful in very few cases.